

Tutorial Week 4

It is the pipeline structure of MIPS CPU meaning that when a jump instruction is performed and a new program counter is generated, the instruction after the jump will be decoded, rather than discarded, the architecture rules state that the instruction after a branch is always executed before the instruction at the target of the branch

Questions

R3000 and assembly

1. What is a *branch delay*?
2. The goal of this question is to have you reverse engineer some of the C compiler function calling convention (instead of reading it from a manual). The following code contains 6 functions that take 1 to 6 integer arguments. Each function sums its arguments and returns the sum as a the result.

```
#include <stdio.h>

/* function prototypes, would normally be in header files */
int arg1(int a);
int arg2(int a, int b);
int arg3(int a, int b, int c);
int arg4(int a, int b, int c, int d);
int arg5(int a, int b, int c, int d, int e );
int arg6(int a, int b, int c, int d, int e, int f);

/* implementations */
int arg1(int a)
{
    return a;
}

int arg2(int a, int b)
{
    return a + b;
}

int arg3(int a, int b, int c)
{
    return a + b + c;
}

int arg4(int a, int b, int c, int d)
{
    return a + b + c + d;
}

int arg5(int a, int b, int c, int d, int e )
{
    return a + b + c + d + e;
}

int arg6(int a, int b, int c, int d, int e, int f)
{
    return a + b + c + d + e + f;
}

/* do nothing main, so we can compile it */
int main()
{
}
```

The following code is the disassembled code that is generated by the C compiler (with certain optimisations turned of for the sake of clarity).

```

004000f0 <arg1>:
  4000f0:      03e00008      jr      ra
  4000f4:      00801021      move   v0,a0

004000f8 <arg2>:
  4000f8:      03e00008      jr      ra
  4000fc:      00851021      addu   v0,a0,a1

00400100 <arg3>:
  400100:      00851021      addu   v0,a0,a1
  400104:      03e00008      jr      ra
  400108:      00461021      addu   v0,v0,a2

0040010c <arg4>:
  40010c:      00852021      addu   a0,a0,a1
  400110:      00861021      addu   v0,a0,a2
  400114:      03e00008      jr      ra
  400118:      00471021      addu   v0,v0,a3

0040011c <arg5>:
  40011c:      00852021      addu   a0,a0,a1
  400120:      00863021      addu   a2,a0,a2
  400124:      00c73821      addu   a3,a2,a3
  400128:      8fa20010      lw     v0,16(sp)
  40012c:      03e00008      jr      ra
  400130:      00e21021      addu   v0,a3,v0

00400134 <arg6>:
  400134:      00852021      addu   a0,a0,a1
  400138:      00863021      addu   a2,a0,a2
  40013c:      00c73821      addu   a3,a2,a3
  400140:      8fa20010      lw     v0,16(sp)
  400144:      00000000      nop
  400148:      00e22021      addu   a0,a3,v0
  40014c:      8fa20014      lw     v0,20(sp)
  400150:      03e00008      jr      ra
  400154:      00821021      addu   v0,a0,v0

00400158 <main>:
  400158:      03e00008      jr      ra
  40015c:      00001021      move   v0,zero

```

because there is no local variable inside the function, so the compiler does not need space on the stack to store them

a. arg1 (and functions in general) returns its return value in what register? **v0**

b. Why is there no stack references in arg2?

c. What does jr ra do? **It jumps to the address in the ra register** because move is in the branch delay, so it is

d. Which register contains the first argument to the function? **a0** executed before arg1 returns

e. Why is the move instruction in arg1 after the jr instruction.

f. Why does arg5 and arg6 reference the stack? **because we only have a0 - a3 in our register, if we need to store more arguments, we need to use stack**

3. The following code provides an example to illustrate stack management by the C compiler. Firstly, examine the C code in the provided example to understand how the recursive function works.

```

#include <stdio.h>
#include <unistd.h>

char teststr[] = "\nThe quick brown fox jumps of the lazy dog.\n";

void reverse_print(char *s)
{
    if (*s != '\0') {
        reverse_print(s+1);
        write(STDOUT_FILENO,s,1);
    }
}

```

```
int main()
{
    reverse_print(teststr);
}
```

The following code is the disassembled code that is generated by the C compiler (with certain optimisations turned off for the sake of clarity).

- Describe what each line in the code is doing.
- What is the maximum depth the stack can grow to when this function is called?

The stack of each invocation of reverse_print is 24 bytes, but the function is recursive. The allocation is 24 bytes times the length of the string, and thus if the string is unbounded, so is the recursion, and thus stack growth is also unbounded.

004000f0 <reverse_print>:	4000f0:	27bdf8e8	addiu	sp, sp, -24	reverse 24 bytes for the stack, 16 for a0-a3, and 8 for ra and s0
	4000f4:	afbf0014	sw	ra, 20(sp)	save the return address for the function on the stack
	4000f8:	afbf0010	sw	s0, 16(sp)	save s0 on the stack so we can use the register in this function,
	4000fc:	80820000	lb	v0, 0(a0)	load the byte at a0 to register v0 as the first argument
	400100:	00000000	nop		v0 won't be ready here
	400104:	10400007	beqz	v0, 400124 <reverse_print+0x34>	if the character (the value of v0) is zero, if so jump to 400124
	400108:	00808021	move	s0, a0	this is on the branch delay slot, save the pointer in s0 to a0
	40010c:	0c10003c	jal	4000f0 <reverse_print>	call reverse_print
	400110:	24840001	addiu	a0, a0, 1	this is at the branch delay, so the value of a0++, move onto next char
	400114:	24040001	li	a0, 1	load the file descriptor for write(1) ???
	400118:	02002821	move	a1, s0	pass the pointer in s0 to write
	40011c:	0c1000af	jal	4002bc <write>	jump and link to 4002bc, call write function
	400120:	24060001	li	a2, 1	this is at the branch delay, load the number of bytes write outputs
	400124:	8fbf0014	lw	ra, 20(sp)	restore the return address of this function in prep for return
	400128:	8fb00010	lw	s0, 16(sp)	restore s0 to whatever it was before this function was called
	40012c:	03e00008	jr	ra	return to the caller
	400130:	27bd0018	addiu	sp, sp, 24	in the branch delay slot, deallocate the stack

- Why is recursion or large arrays of local variables avoided by kernel programmers?


because we have limited size of kernel stack and it can grow very fast, we don't want stack overflow crashes the entire machine

Threads

in the cooperative case you have to manually yield, so you don't take any resources from other parts of the system.

but for preemptive case, we don't have to worry about it because the OS handles everything for us

- Compare cooperative versus preemptive multithreading?



- Describe *user-level threads* and *kernel-level threads*. What are the advantages or disadvantages of each approach?

they run in user space, kernel level threads are managed by the kernel but not ran in the kernel

user-level threads: can be run on an os that doesn't support kernel threads; faster to create, destroy etc

kernel-level threads take advantage of multi-core cpus usually preemptive slow and inefficient, requires kernel entry and exit, more overhead

usually cooperative, some user level threads use alarms or timeouts to provide a tick for preemption blocking syscalls freeze up the entire process

- A web server is constructed such that it is multithreaded. If the only way to read from a file is a normal blocking read system call, do you think user-level threads or kernel-level threads are being used for the web server? Why?

kernel level threads are being used because if the server is constructed using user-level threads, when someone is performing a read, other people will be blocked from the process until the user has finished reading, or syscall returns

- Assume a multi-process operating system with single-threaded applications. The OS manages the concurrent application requests by having a *thread* of control within the kernel for each process. Such a OS would have an in-kernel stack associated with each process.

Switching between each process (in-kernel thread) is performed by the function `switch_thread(cur_tcb, dst_tcb)`. What does this function do?

Kernel Entry and Exit

- What is the EPC register? What is it used for?

This is a 32-bit register containing the 32-bit address of the return point for the last exception. The instruction causing (or suffering) the exception is at EPC, unless BD is set in Cause, in which case EPC points to the previous (branch) instruction. It is used by the exception handler to restart execution at the point where execution was interrupted.

The 'c' (current) bits are shifted into the corresponding 'p' (previous) bits, after which KUC = 0, IEC = 0 (kernel mode with interrupts disabled). They are shifted in order to preserve the current state at the point of the exception in order to restore that exact state when returning from the exception. They are restored via a rfe instruction (restore from exception).

10. What happens to the KUC and IEC bits in the STATUS register when an exception occurs? Why? How are they restored?

11. What is the value of ExcCode in the Cause register immediately after a system call exception occurs?
the value will be corresponding to the number of the syscall exception - 8

12. Why must kernel programmers be especially careful when implementing system calls?

System calls with poor argument checking or implementation can result in a malicious or buggy program crashing, or compromising the operating system.

v0 register contains what type of syscalls. a0-a3 are the arguments

13. The following questions are focused on the case study of the system call convention used by OS/161 on the MIPS R3000 from the lecture slides.

1. How does the 'C' function calling convention relate to the system-call interface between the application and the kernel?

The 'C' function calling convention is what library code has to be adhered to after any system-call wrapper function completes. This involves saving and restoring of the preserved registers (e.g., s0-s6, ra). The system call convention also uses the calling convention of the C-compiler to pass arguments to OS/161. Having the same convention as the compiler means the system call wrapper can avoid moving arguments around and the compiler has already placed them where the OS expects to find them.

2. What does the most work to preserve the compiler calling convention, the system call wrapper, or the OS/161 kernel.

The OS/161 kernel code does the saving and restoring of preserved registers. The system call wrapper function does very little.

3. At minimum, what additional information is required beyond that passed to the system-call wrapper function?

The interface between the system-call wrapper function and the kernel can be defined to provide additional information beyond that passed to the wrapper function. At minimum, the wrapper function must add the system call number to the arguments passed to the wrapper function. It's usually added by setting an agreed-to register to the value of the system call number.

14. In the example given in lectures, the library function *read* invoked the read system call. Is it essential that both have the same name? If not, which name is important?

system calls don't really have names, the name is there just for us to understand what system calls are invoked. The library function *read()* traps in to the kernel, it puts the number of the system call into a register or on the stack. This number is used to index into a jump table. There is really no name used anywhere. On the other hand, the name of the library function is very important as it is what appears in the program.

15. To a programmer, a system call looks like any other call to a library function. Is it important that a programmer know which library function result in system calls? Under what circumstances and why?

In performance critical, syscalls are not ideal because we need context switch between user space and kernel space, which is a performance overhead.

As far as program logic is concerned it does not matter whether a call to a library function results in a system call. But if performance is an issue, if a task can be accomplished without a system call the program will run faster.

16. Describe a plausible sequence of activities that occur when a timer interrupt results in a context switch.

Every system call involves overhead in switching from the user context to the kernel context. Furthermore, on a multiuser system the operating system may schedule another process to run when a system call completes, further slowing the progress in real time of a calling process.